DeepSynth: Large-Scale Training (Test Paper 4)

SYNTHETIC TEST PAPER - hypothetical values created for testing ANVIKSHA. Not real experimental findings.

Results

A ResNet-50 trained for 200 epochs on CIFAR-100 reaches 94.1% top-1 accuracy.

Inference is 3x faster than the baseline on a single A100 GPU.

We believe the learned features are more interpretable than prior work.